

GrabCut and Semantic Segmentation

Tejas Sharma, Jiajun Xu, Youxin Chen, Jacob Irle

December 14, 2017

Introduction

The goal of our project was to implement Grabcut, which is a foreground extraction algorithm, and then to further utilize this tool in Semantic Segmentation, which, when trained with a set of pre-labeled images, can complete GrabCut without any human interaction at all for that class of objects.

1 GrabCut

1.1 Outside Resources

- Our grabcut implementation is based on the paper “ ‘GrabCut’ -- Interactive Foreground Extraction using Iterated Graph Cuts” by Carsten Rother, Vladimir Kolmogorov, and Andrew Blake.
- GrabCut is based upon the paper “Graph Cut”, by Yuri Boykov and Marie-Pierre Jolly.
- The paper “Implementing GrabCut” by Justin Talbot and Xiaoqian Xu was also crucial in our implementation, as it explained the method in a much simpler manner, which increased our understanding significantly.
- We also used the “PyMaxFlow” library, which is a library for solving min-cut/max-flow in python. This library is based on the method described in “An Experimental Comparison of Min-Cut/Max-Flow Algorithms for Energy Minimization in Vision” by Yuri Boykov and Vladimir Kolmogorov. We initially created our own implementation based on the same paper, but found that it was too slow to be practically used due to being written in pure python. PyMaxFlow creates a wrapper for the C++ version, which makes it a lot faster.
- Our UI is heavily based upon the Python example available in the OpenCV github repo.

1.2 Implementation Details

GrabCut is based upon the Graph Cut algorithm, which is also used for image segmentation. The key difference between the two is that Graph Cut could only be used in black and white images, and required a lot more user interaction.

It uses iterative learning, which greatly reduces the amount of user interaction, but still allows the user to hardcode foreground and background pixels. At a high level, GrabCut attempts to minimize the energy between pixels in an image in order to find an optimal bounding cut. It iteratively improves this cut by recalculating the energy based on its previous classification until convergence.

1.2.1 Gaussian Mixture Models

GrabCut represents the foreground and the background of an image with 2 Gaussian Mixture Models, each of which has K gaussian components. Every pixel of the image belongs to one of the foreground or background component. Each gaussian component stores the mean (μ), the covariance matrix (Σ) and a weight (π). We use these gaussian components to calculate the likelihood that the pixel belongs to either the foreground or the background

We calculate the likelihood a pixel belongs to a GMM using equation 2 from “Implementing GrabCut”, which gives us a likelihood as the function of a pixel m

$$D(m) = -\log \sum_{i=1}^K \left[\pi(\alpha_m, i) \frac{1}{\sqrt{\det \Sigma(\alpha_m, i)}} \times \exp \left(\frac{1}{2} [z_m - \mu(\alpha_m, i)]^T \Sigma(\alpha_m, i)^{-1} [z_m - \mu(\alpha_m, i)] \right) \right]$$

where z_m is the colour of the pixel, and (α_m, i) essentially allows us to uniquely identify the cluster covariance and mean that we are trying to access.

1.2.2 Pixel Classifications

When creating the GMMs, we use the initial user input, which is a simple rectangle that is drawn around the foreground object, to create a set of classifications for the pixels. In “Implementing GrabCut”, they use a trimap of Foreground, Background and Unknown. However, when looking at the example code for our UI, we noticed that it uses four values: Background, Foreground, Potential Background, and Potential Foreground. We decided this made more sense to us, and would make implementation easier, so we also used this method. In the initial step, everything outside the rectangle is classified as Background, and everything inside it is classified as Potential Foreground. To note the difference in this report between the classification Foreground and the regular word foreground, we will be capitalising the classifications.

1.2.3 Graph Cut

To create the graph used in Graph Cut, we represent each pixel as a node in the graph. Each pixel is connected to its 8 adjacent pixels in the image (unless it is on one or more of the borders). We also add two terminal nodes for performing the max flow segmentation, one representing the source (background), and one representing the sink (foreground). The terminal nodes are connected to every pixel.

There are two types of edges in the graph: T-links, which are between a pixel and the two terminal nodes; and N-links, which are between a pixel and its 8 neighbours.

We build the N-links as follows:

- Each pixel is connected to its neighbours with the weight function

$$N(m, n) = \frac{1}{\text{dist}(m, n)} e^{-\beta \|z_m - z_n\|^2}$$

- $\beta = \left(2 \langle (z_m - z_n)^2 \rangle\right)^{-1}$, where $\langle \rangle$ represents an expectation over an image sample.
- The *dist* factor simply denotes whether the pixels are diagonal neighbours or immediate neighbours as noted in “Implementing GrabCut”, since the method in the original GrabCut paper favoured diagonal edges.

We build the T-links between the terminal nodes and the pixels by doing the following:

- First, we store the maximum weight created in the N-links. This will be used when a pixel has been classified as Foreground or Background.
- If a pixel is classified as Foreground, we assign its foreground T-link to be the maximum weight, and its background T-link to be 0
- If a pixel is classified as Background, we assign its foreground T-link to be 0, and its background T-link to be the maximum weight
- If a pixel is classified as Potential Background or Potential Foreground, we use the $D(m)$ function referenced earlier and assign its foreground T-link to be $D_{bgd}(m)$ and its background T-link to be $D_{fgd}(m)$.

After building the links, we run a max cut algorithm on it, and use the results to tell us which parts are foreground and background. Pixels that have been marked as Foreground or Background by the user are unchanged, and pixels that are not are updated to Potential Foreground and Potential Background based on the cut. These values are then fed back into the GMMs to iteratively run until convergence.

1.3 GrabCut Results

The following are the results of running our GrabCut implementation on some images that were used in the original papers.

The first is an image of a llama, which we decided worked pretty well without any secondary user input.

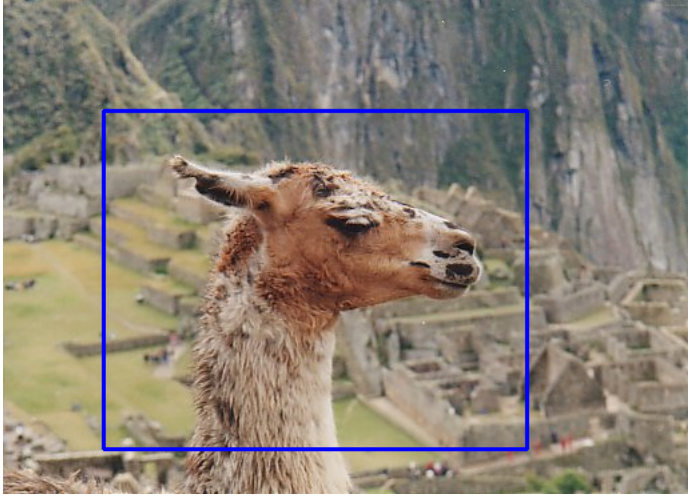
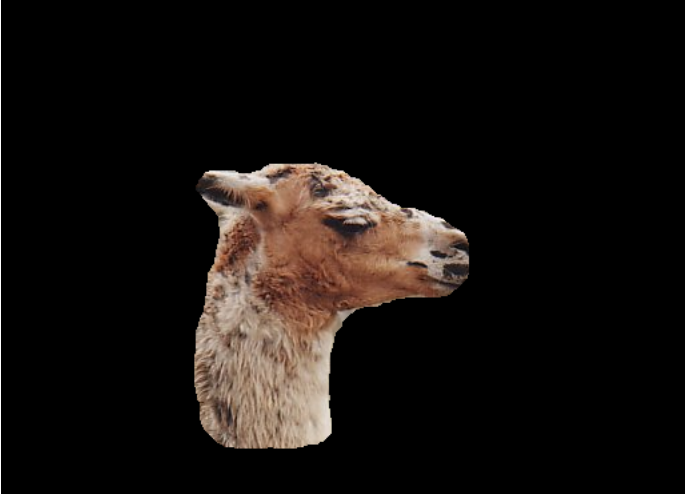
Original Image	GrabCut Output
 The original image shows a brown llama in profile, facing right, in the foreground. In the background, the ancient stone ruins of Machu Picchu are visible, nestled in a valley with steep, green mountains in the distance. A blue rectangular bounding box is drawn around the llama's head and neck.	 The GrabCut segmentation output shows the llama from the original image isolated against a solid black background. The segmentation is precise, following the contours of the llama's head, neck, and body.

Table 1: Llama Segmentation

The next is an image of Lena Söderberg, which allowed us to show off the ability for users to mark certain areas as Foreground or Background that were incorrectly classified by GrabCut.



Original Image	Initial GrabCut Output
	
User Touchup	Final GrabCut Output
	

Table 2: Lena Segmentation

Finally, we have an image of a soldier. We were expecting this image to be really hard due to the similarity of colours between the foreground and background. Indeed, the initial result was not great. However, even with such similar colours, after some user input we had a pretty good result.





Original Image	Initial GrabCut Output
	
User Touchup	Final GrabCut Output
	

Table 3: Lena Segmentation

2 Semantic Segmentation

2.1 Outside Resources

- Our Semantic Segmentation is based on “Semantic Segmentation using GrabCut” by Christopher Göring, Björn Fröhlich, and Joachim Denzler.
- The metric we use to decide the similarity between two patches is Hu Moments, based on the paper “Visual Pattern Recognition by Moment Invariants” by Ming-Kuei Hu.

2.2 Implementation Details